

# tomer-churn-prediction-using-ml-1

December 18, 2025

## 0.1 1.Importing the dependencies

```
[104]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,classification_report
import pickle
```

## 0.2 2. Data Loading and Understanding

```
[107]: # load thecsv data to a pandas dataframe
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
```

```
[109]: df.shape
```

```
[109]: (7043, 21)
```

```
[110]: df.head()
```

```
[110]:    customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService \
0    7590-VHVEG  Female           0      Yes        No         1        No
1    5575-GNVDE    Male           0       No        No        34      Yes
2    3668-QPYBK    Male           0       No        No         2      Yes
3    7795-CFOCW    Male           0       No        No        45        No
4    9237-HQITU  Female           0       No        No         2      Yes

          MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0  No phone service            DSL            No          Yes
1                  No            DSL           Yes          No
```

```

2           No          DSL        Yes       Yes
3  No phone service      DSL        Yes       No
4           No     Fiber optic    No        No

DeviceProtection TechSupport StreamingTV StreamingMovies      Contract \
0           No          No         No        No Month-to-month
1          Yes         No         No        No One year
2           No          No         No        No Month-to-month
3          Yes         Yes        No        No One year
4           No          No         No        No Month-to-month

PaperlessBilling      PaymentMethod MonthlyCharges TotalCharges \
0           Yes        Electronic check   29.85      29.85
1           No         Mailed check    56.95     1889.5
2          Yes        Mailed check    53.85      108.15
3           No     Bank transfer (automatic) 42.30     1840.75
4           Yes        Electronic check  70.70      151.65

Churn
0     No
1     No
2    Yes
3     No
4    Yes

```

[111]: df.tail()

```

[111]:   customerID  gender SeniorCitizen Partner Dependents tenure \
7038  6840-RESVB   Male        0     Yes     Yes    24
7039  2234-XADUH Female       0     Yes     Yes    72
7040  4801-JZAZL Female       0     Yes     Yes    11
7041  8361-LTMKD   Male        1     Yes     No     4
7042  3186-AJIEK   Male        0     No      No    66

PhoneService      MultipleLines InternetService OnlineSecurity \
7038          Yes            Yes          DSL        Yes
7039          Yes            Yes     Fiber optic     No
7040          No  No phone service      DSL        Yes
7041          Yes            Yes     Fiber optic     No
7042          Yes            No     Fiber optic        Yes

OnlineBackup DeviceProtection TechSupport StreamingTV StreamingMovies \
7038          No            Yes          Yes        Yes       Yes
7039          Yes           Yes          No        Yes       Yes
7040          No            No           No        No        No
7041          No            No           No        No        No
7042          No            Yes          Yes        Yes       Yes

```

	Contract	PaperlessBilling	PaymentMethod	\
7038	One year	Yes	Mailed check	
7039	One year	Yes	Credit card (automatic)	
7040	Month-to-month	Yes	Electronic check	
7041	Month-to-month	Yes	Mailed check	
7042	Two year	Yes	Bank transfer (automatic)	

	MonthlyCharges	TotalCharges	Churn	
7038	84.80	1990.5	No	
7039	103.20	7362.9	No	
7040	29.60	346.45	No	
7041	74.40	306.6	Yes	
7042	105.65	6844.5	No	

```
[112]: pd.set_option("display.max_columns", None)
```

```
[113]: df.head(2)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	No phone service	DSL	No	Yes	
1	No	DSL	Yes	No	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn	
0	Yes	Electronic check	29.85	29.85	No	
1	No	Mailed check	56.95	1889.5	No	

```
[114]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 ...   ...             ...           ...    
```

```

5    tenure            7043 non-null   int64
6    PhoneService      7043 non-null   object
7    MultipleLines     7043 non-null   object
8    InternetService   7043 non-null   object
9    OnlineSecurity    7043 non-null   object
10   OnlineBackup       7043 non-null   object
11   DeviceProtection  7043 non-null   object
12   TechSupport        7043 non-null   object
13   StreamingTV        7043 non-null   object
14   StreamingMovies    7043 non-null   object
15   Contract           7043 non-null   object
16   PaperlessBilling   7043 non-null   object
17   PaymentMethod       7043 non-null   object
18   MonthlyCharges     7043 non-null   float64
19   TotalCharges        7043 non-null   object
20   Churn              7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
[115]: # dropping customerID column as this is not required for modelling
df = df.drop(columns=["customerID"])
```

```
[116]: df.head(2)
```

```

[116]: gender SeniorCitizen Partner Dependents tenure PhoneService \
0 Female          0 Yes No 1 No
1 Male            0 No No 34 Yes

MultipleLines InternetService OnlineSecurity OnlineBackup \
0 No phone service DSL No Yes
1 No DSL Yes No

DeviceProtection TechSupport StreamingTV StreamingMovies Contract \
0 No No No No Month-to-month
1 Yes No No No One year

PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn
0 Yes Electronic check 29.85 29.85 No
1 No Mailed check 56.95 1889.5 No

```

```
[117]: df.columns
```

```
[117]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
       'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
```

```
        dtype='object')

[118]: print(df["gender"].unique())

['Female' 'Male']

[119]: print(df["SeniorCitizen"].unique())

[0 1]

[120]: # printing the unique values in all the columns

numerical_features_list = ["tenure", "MonthlyCharges", "TotalCharges"]

for col in df.columns:
    if col not in numerical_features_list:
        print(col, df[col].unique())
        print("-" * 50)

gender ['Female' 'Male']
-----
SeniorCitizen [0 1]
-----
Partner ['Yes' 'No']
-----
Dependents ['No' 'Yes']
-----
PhoneService ['No' 'Yes']
-----
MultipleLines ['No phone service' 'No' 'Yes']
-----
InternetService ['DSL' 'Fiber optic' 'No']
-----
OnlineSecurity ['No' 'Yes' 'No internet service']
-----
OnlineBackup ['Yes' 'No' 'No internet service']
-----
DeviceProtection ['No' 'Yes' 'No internet service']
-----
TechSupport ['No' 'Yes' 'No internet service']
-----
StreamingTV ['No' 'Yes' 'No internet service']
-----
StreamingMovies ['No' 'Yes' 'No internet service']
-----
Contract ['Month-to-month' 'One year' 'Two year']
-----
PaperlessBilling ['Yes' 'No']
```

```
-----  
PaymentMethod ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'  
 'Credit card (automatic)']  
-----
```

```
Churn ['No' 'Yes']  
-----
```

```
[121]: print(df.isnull().sum())
```

```
gender          0  
SeniorCitizen   0  
Partner         0  
Dependents     0  
tenure          0  
PhoneService    0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport     0  
StreamingTV    0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
TotalCharges    0  
Churn           0  
dtype: int64
```

```
[122]: #df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
[123]: df[df["TotalCharges"]== " "]
```

```
[123]:      gender SeniorCitizen Partner Dependents tenure PhoneService \
488   Female          0     Yes      Yes      0        No
753    Male          0      No      Yes      0       Yes
936   Female          0     Yes      Yes      0       Yes
1082   Male          0     Yes      Yes      0       Yes
1340   Female          0     Yes      Yes      0        No
3331   Male          0     Yes      Yes      0       Yes
3826   Male          0     Yes      Yes      0       Yes
4380   Female          0     Yes      Yes      0       Yes
5218   Male          0     Yes      Yes      0       Yes
6670   Female          0     Yes      Yes      0       Yes
6754   Male          0      No      Yes      0       Yes
```

	MultipleLines	InternetService	OnlineSecurity	\	
488	No phone service	DSL	Yes		
753	No	No	No internet service		
936	No	DSL	Yes		
1082	Yes	No	No internet service		
1340	No phone service	DSL	Yes		
3331	No	No	No internet service		
3826	Yes	No	No internet service		
4380	No	No	No internet service		
5218	No	No	No internet service		
6670	Yes	DSL	No		
6754	Yes	DSL	Yes		
	OnlineBackup	DeviceProtection	TechSupport	\	
488	No	Yes	Yes		
753	No internet service	No internet service	No internet service		
936	Yes	Yes	No		
1082	No internet service	No internet service	No internet service		
1340	Yes	Yes	Yes		
3331	No internet service	No internet service	No internet service		
3826	No internet service	No internet service	No internet service		
4380	No internet service	No internet service	No internet service		
5218	No internet service	No internet service	No internet service		
6670	Yes	Yes	Yes		
6754	Yes	No	Yes		
	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
488	Yes	No	Two year	Yes	
753	No internet service	No internet service	Two year	No	
936	Yes	Yes	Two year	No	
1082	No internet service	No internet service	Two year	No	
1340	Yes	No	Two year	No	
3331	No internet service	No internet service	Two year	No	
3826	No internet service	No internet service	Two year	No	
4380	No internet service	No internet service	Two year	No	
5218	No internet service	No internet service	One year	Yes	
6670	Yes	No	Two year	No	
6754	No	No	Two year	Yes	
	PaymentMethod	MonthlyCharges	TotalCharges	Churn	
488	Bank transfer (automatic)	52.55		No	
753	Mailed check	20.25		No	
936	Mailed check	80.85		No	
1082	Mailed check	25.75		No	
1340	Credit card (automatic)	56.05		No	
3331	Mailed check	19.85		No	

3826	Mailed check	25.35	No
4380	Mailed check	20.00	No
5218	Mailed check	19.70	No
6670	Mailed check	73.35	No
6754	Bank transfer (automatic)	61.90	No

```
[124]: len(df[df["TotalCharges"]==" "])
```

```
[124]: 11
```

```
[125]: df["TotalCharges"] = df["TotalCharges"].replace({" ":" "0.0"})
```

```
[126]: df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
[127]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   gender          7043 non-null   object 
 1   SeniorCitizen   7043 non-null   int64  
 2   Partner         7043 non-null   object 
 3   Dependents     7043 non-null   object 
 4   tenure          7043 non-null   int64  
 5   PhoneService    7043 non-null   object 
 6   MultipleLines   7043 non-null   object 
 7   InternetService 7043 non-null   object 
 8   OnlineSecurity  7043 non-null   object 
 9   OnlineBackup    7043 non-null   object 
 10  DeviceProtection 7043 non-null   object 
 11  TechSupport    7043 non-null   object 
 12  StreamingTV    7043 non-null   object 
 13  StreamingMovies 7043 non-null   object 
 14  Contract        7043 non-null   object 
 15  PaperlessBilling 7043 non-null   object 
 16  PaymentMethod   7043 non-null   object 
 17  MonthlyCharges  7043 non-null   float64 
 18  TotalCharges    7043 non-null   float64 
 19  Churn           7043 non-null   object 
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

```
[128]: # checking the class distribution of target column
print(df["Churn"].value_counts())
```

Churn

```
No      5174  
Yes     1869  
Name: count, dtype: int64
```

### 0.2.1 Insights:

0.2.2 Customer ID removed as it is not required for modelling

0.2.3 No mmissing values in the dataset

0.2.4 Missing values in the TotalCharges column were replaced with 0

0.2.5 Class imbalance identified in the target

## 0.3 3.. Exploratory Data Analysis (EDA)

```
[129]: df.shape
```

```
[129]: (7043, 20)
```

```
[130]: df.columns
```

```
[130]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',  
           'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',  
           'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',  
           'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',  
           'MonthlyCharges', 'TotalCharges', 'Churn'],  
           dtype='object')
```

```
[131]: df.head(2)
```

```
[131]: gender  SeniorCitizen Partner Dependents  tenure PhoneService  \  
0  Female          0      Yes        No       1        No  
1    Male          0      No        No      34        Yes  
  
      MultipleLines InternetService OnlineSecurity OnlineBackup  \  
0  No phone service          DSL        No        Yes  
1            No          DSL        Yes        No  
  
      DeviceProtection TechSupport StreamingTV StreamingMovies      Contract  \  
0            No          No        No        No Month-to-month  
1            Yes          No        No        No One year  
  
      PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn  
0            Yes  Electronic check        29.85      29.85      No  
1            No    Mailed check        56.95  1889.50      No
```

```
[133]: df.describe()
```

```
[133]:      SeniorCitizen      tenure  MonthlyCharges  TotalCharges
count    7043.000000  7043.000000    7043.000000  7043.000000
mean     0.162147    32.371149     64.761692  2279.734304
std      0.368612    24.559481     30.090047  2266.794470
min      0.000000    0.000000     18.250000     0.000000
25%     0.000000    9.000000     35.500000   398.550000
50%     0.000000   29.000000     70.350000  1394.550000
75%     0.000000   55.000000     89.850000  3786.600000
max     1.000000   72.000000    118.750000  8684.800000
```

### 0.3.1 Numerical Features - Analysis

### 0.3.2 Understand the distribution of teh numerical features

```
[135]: def plot_histogram(df, column_name):

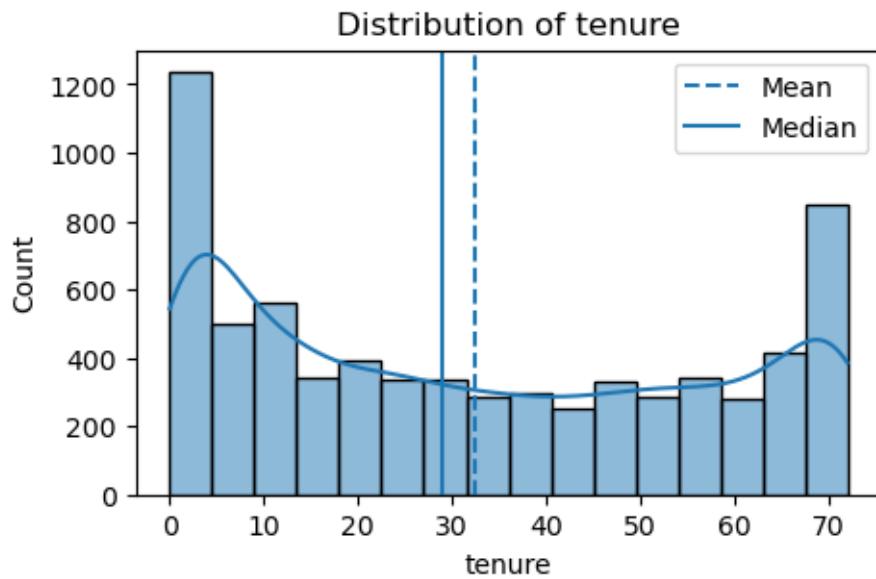
    plt.figure(figsize=(5, 3))
    sns.histplot(df[column_name], kde=True)
    plt.title(f"Distribution of {column_name}")

    # calculate mean and median
    col_mean = df[column_name].mean()
    col_median = df[column_name].median()

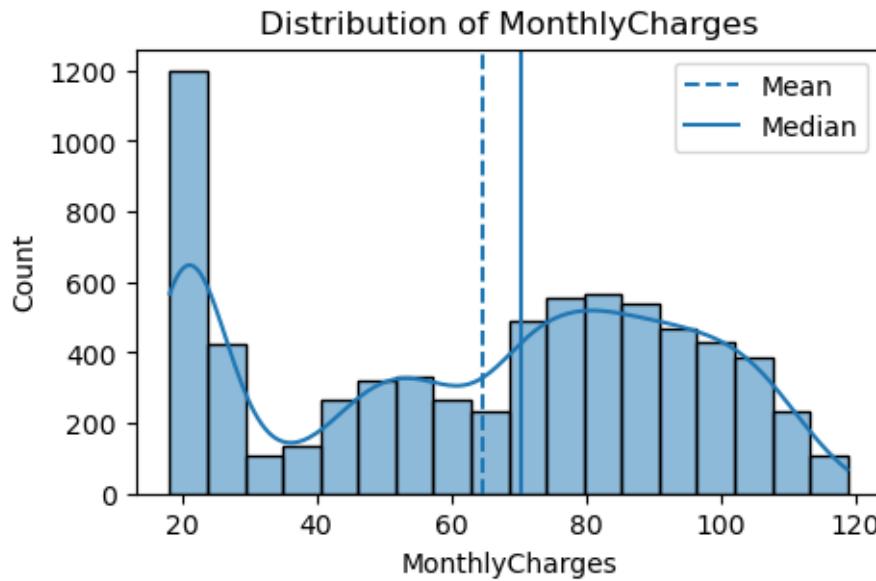
    # add vertical lines
    plt.axvline(col_mean, linestyle="--", label="Mean")
    plt.axvline(col_median, linestyle="-", label="Median")

    plt.legend()
    plt.show()
```

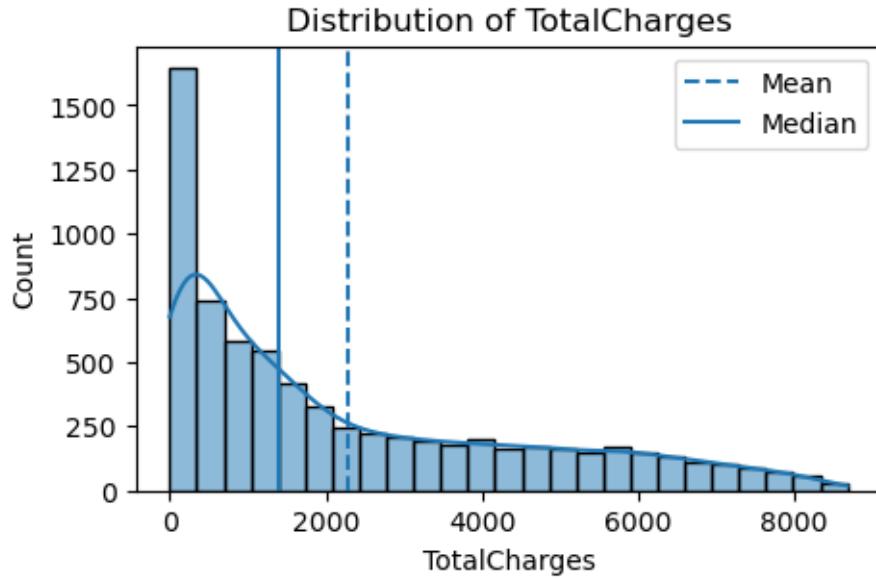
```
[136]: plot_histogram(df, "tenure")
```



```
[137]: plot_histogram(df, "MonthlyCharges")
```



```
[138]: plot_histogram(df, "TotalCharges")
```



#### 0.3.3 Box plot for numerical features

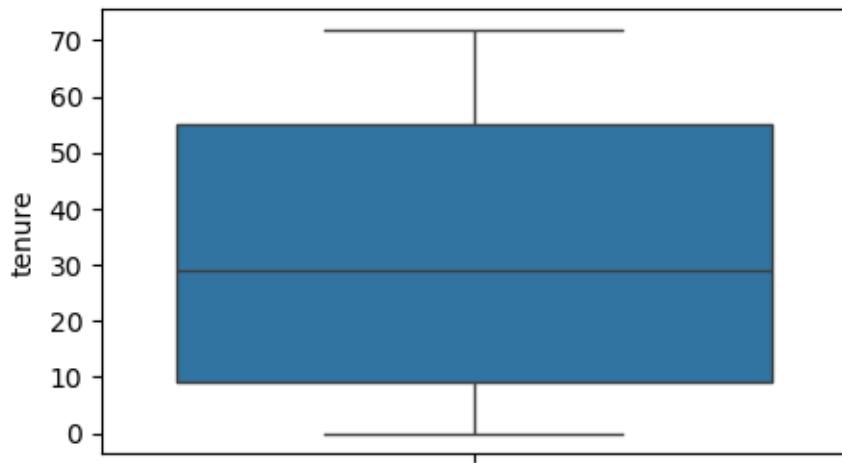
```
[141]: import matplotlib.pyplot as plt

def plot_boxplot(df, column_name):

    plt.figure(figsize=(5, 3))
    sns.boxplot(y=df[column_name])
    plt.title(f"Box Plot of {column_name}")
    plt.ylabel(column_name)
    plt.show()
```

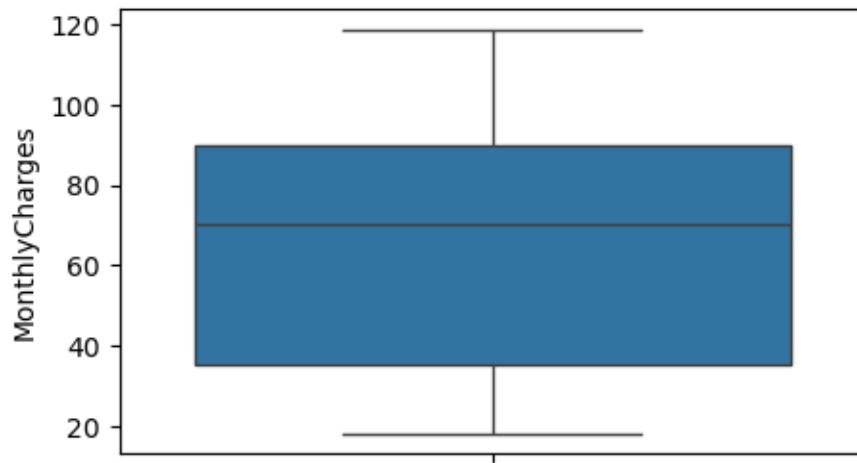
```
[142]: plot_boxplot(df, "tenure")
```

Box Plot of tenure

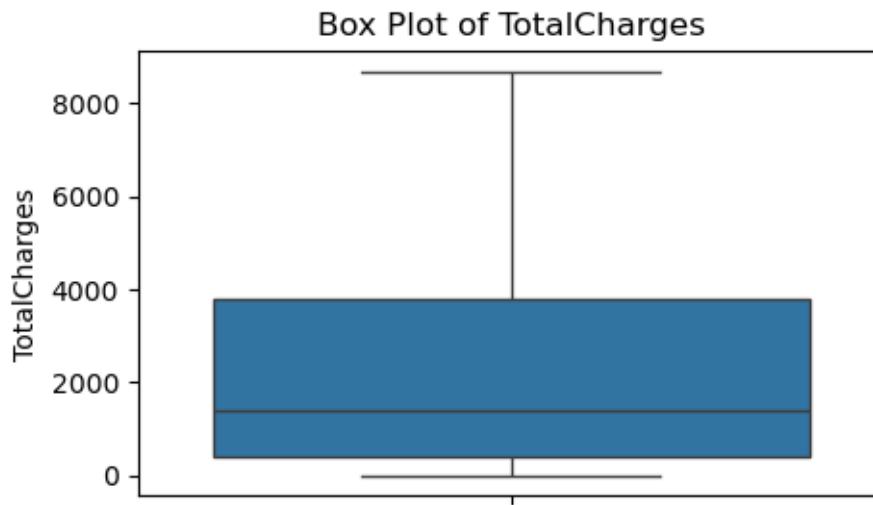


```
[143]: plot_boxplot(df, "MonthlyCharges")
```

Box Plot of MonthlyCharges

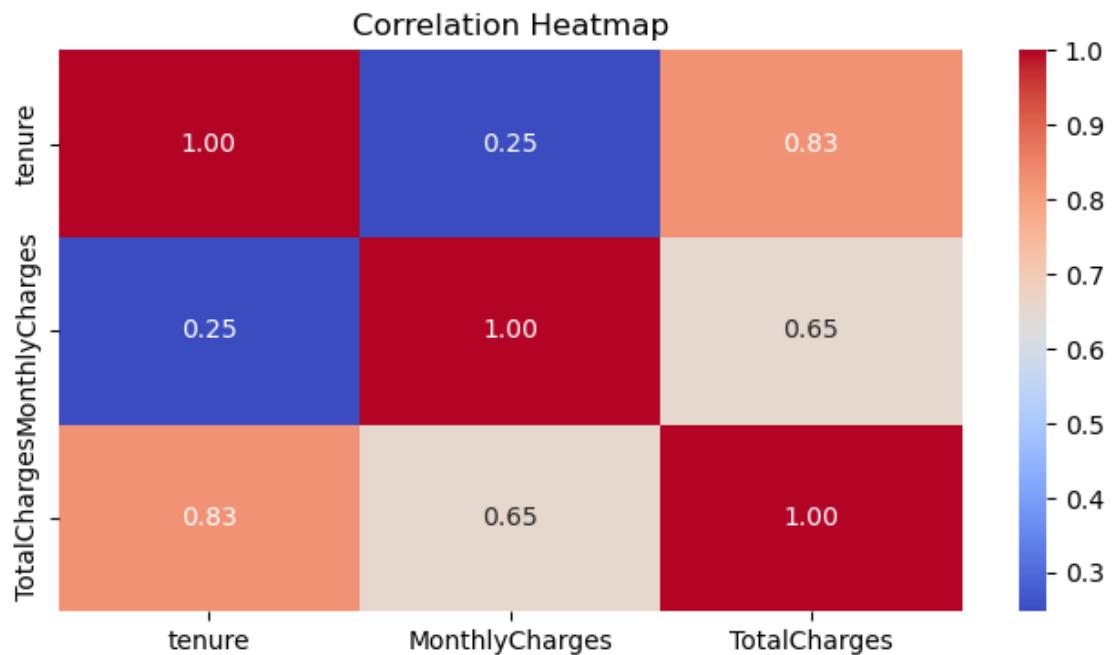


```
[144]: plot_boxplot(df, "TotalCharges")
```



### 0.3.4 Correlation Heatmap for numerical columns

```
[145]: # correlation matrix - heatmap
plt.figure(figsize=(8, 4))
sns.heatmap(df[["tenure", "MonthlyCharges", "TotalCharges"]].corr(), □
    annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



### 0.3.5 Categorical features - Analysis

```
[147]: df.columns
```

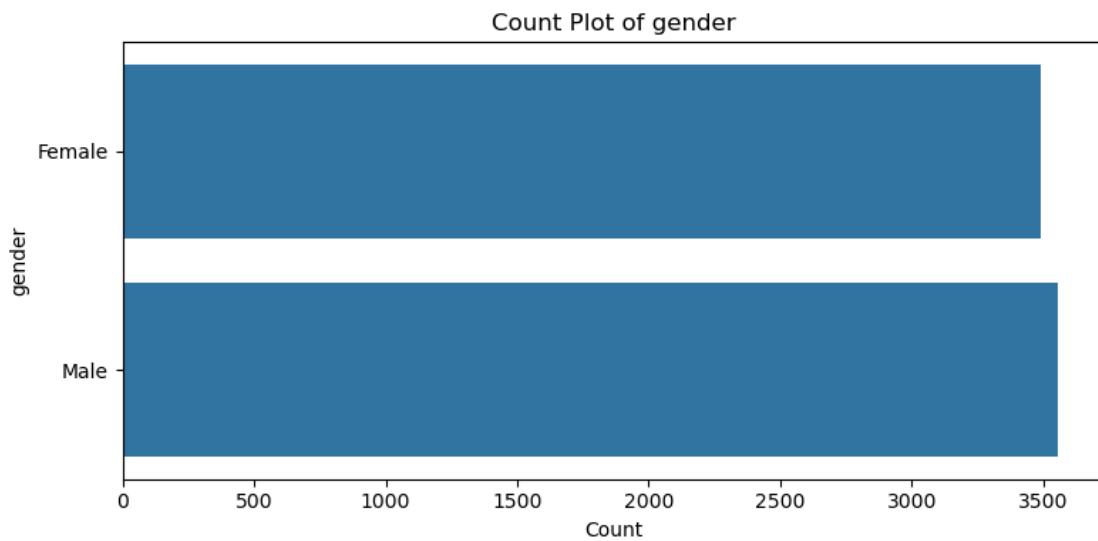
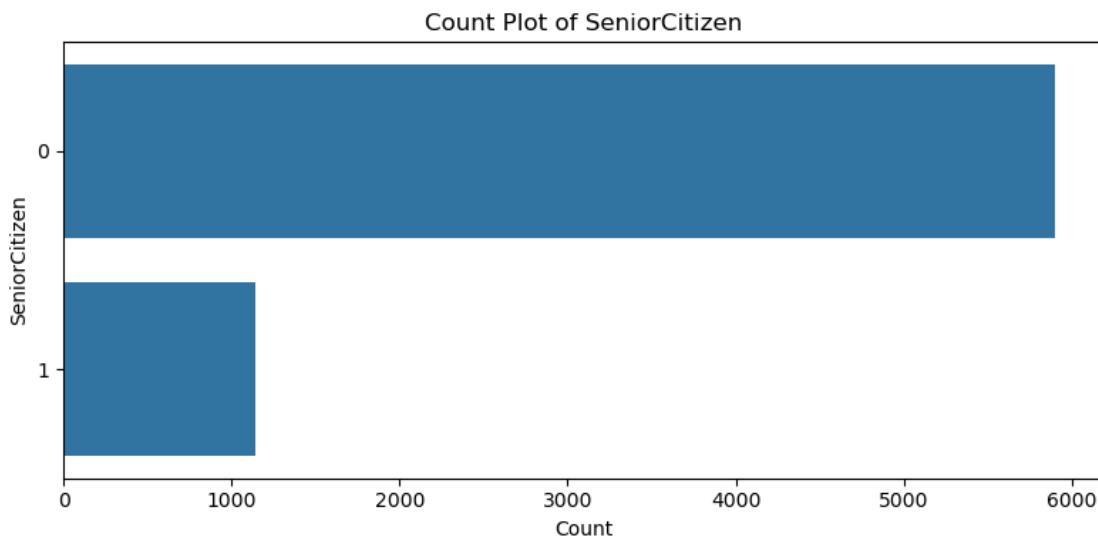
```
[147]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
   'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
   'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
   'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
   'MonthlyCharges', 'TotalCharges', 'Churn'],
  dtype='object')
```

```
[148]: df.info()
```

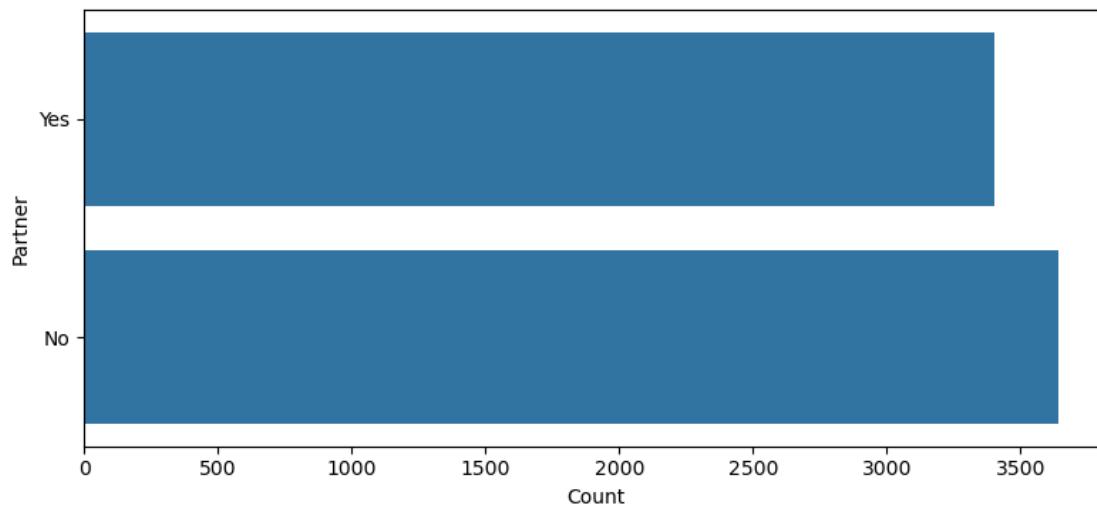
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #  Column          Non-Null Count  Dtype  
---  --  
 0   gender          7043 non-null   object 
 1   SeniorCitizen   7043 non-null   int64  
 2   Partner          7043 non-null   object 
 3   Dependents      7043 non-null   object 
 4   tenure           7043 non-null   int64  
 5   PhoneService     7043 non-null   object 
 6   MultipleLines    7043 non-null   object 
 7   InternetService  7043 non-null   object 
 8   OnlineSecurity   7043 non-null   object 
 9   OnlineBackup      7043 non-null   object 
 10  DeviceProtection 7043 non-null   object 
 11  TechSupport      7043 non-null   object 
 12  StreamingTV       7043 non-null   object 
 13  StreamingMovies   7043 non-null   object 
 14  Contract          7043 non-null   object 
 15  PaperlessBilling  7043 non-null   object 
 16  PaymentMethod     7043 non-null   object 
 17  MonthlyCharges    7043 non-null   float64
 18  TotalCharges      7043 non-null   float64
 19  Churn             7043 non-null   object 
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

### 0.3.6 Countplot for categorical columns

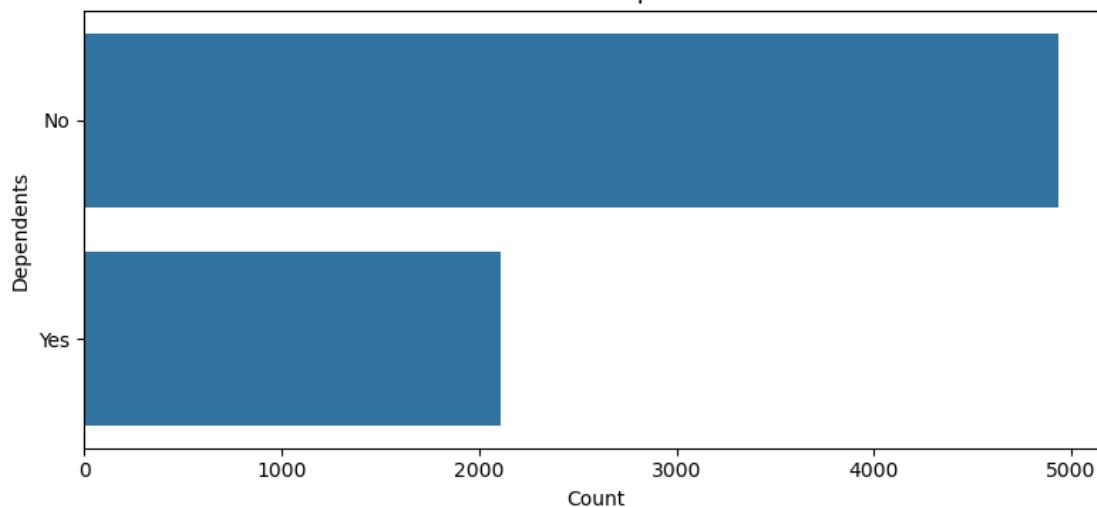
```
[151]: for col in object_cols:  
    plt.figure(figsize=(8, 4))  
    sns.countplot(y=df[col])  
    plt.title(f"Count Plot of {col}")  
    plt.xlabel("Count")  
    plt.ylabel(col)  
    plt.tight_layout()  
    plt.show()
```



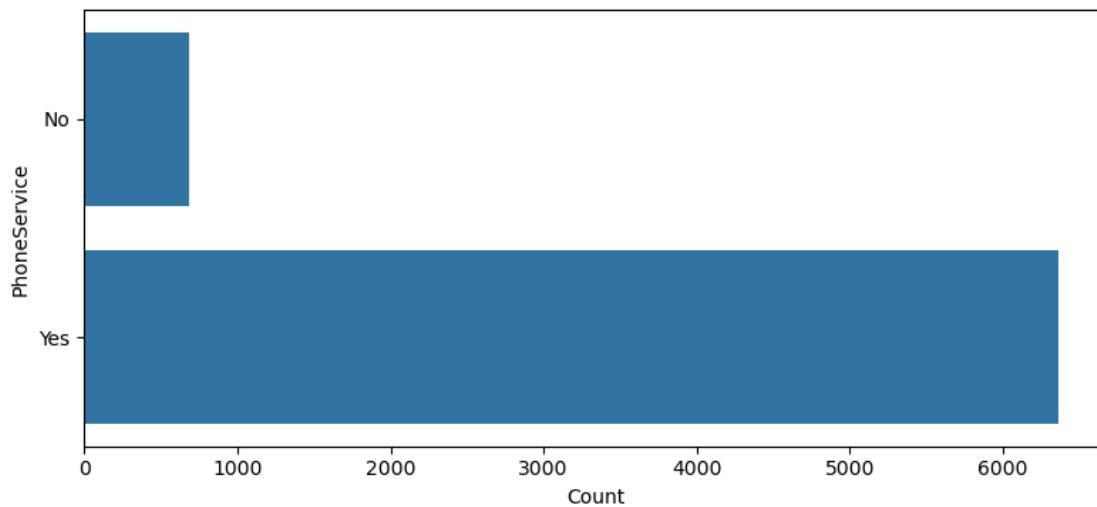
Count Plot of Partner



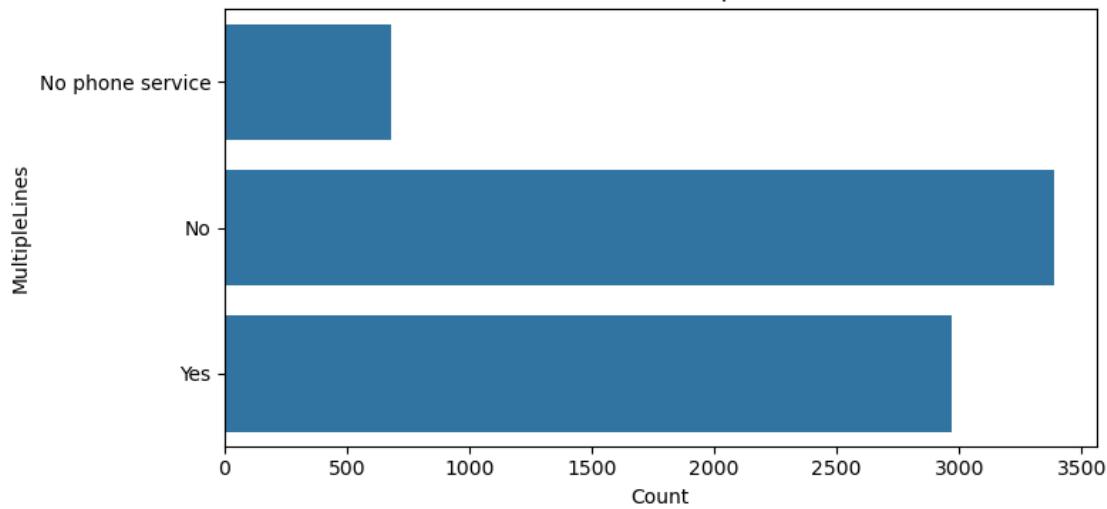
Count Plot of Dependents

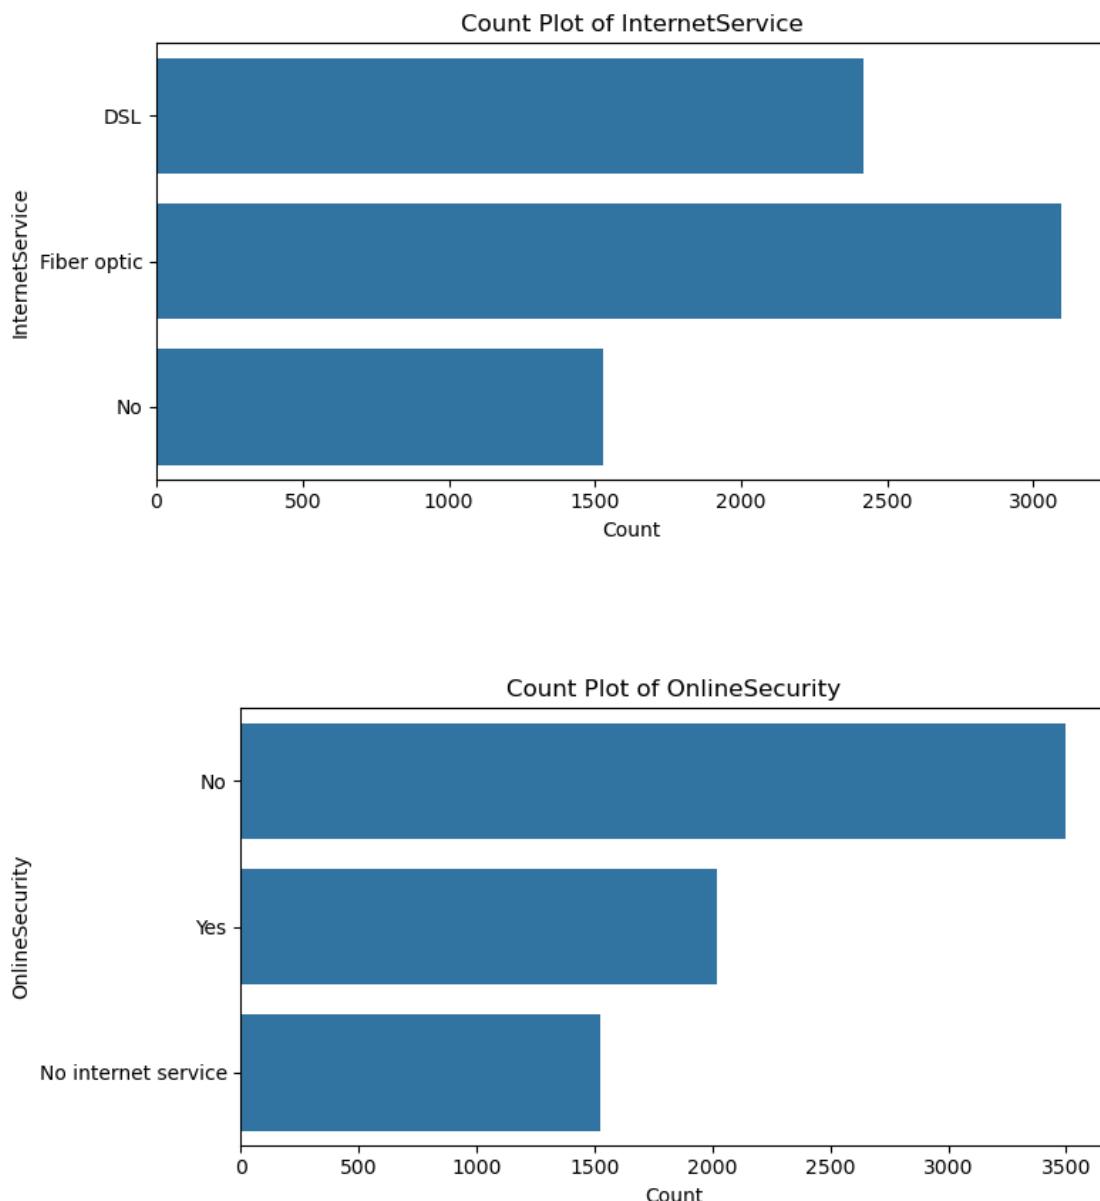


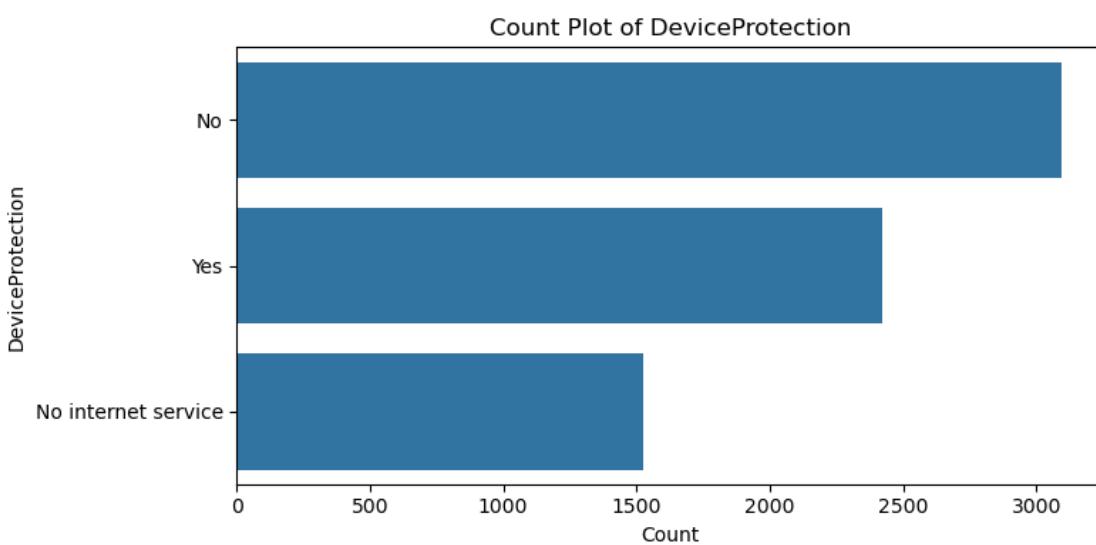
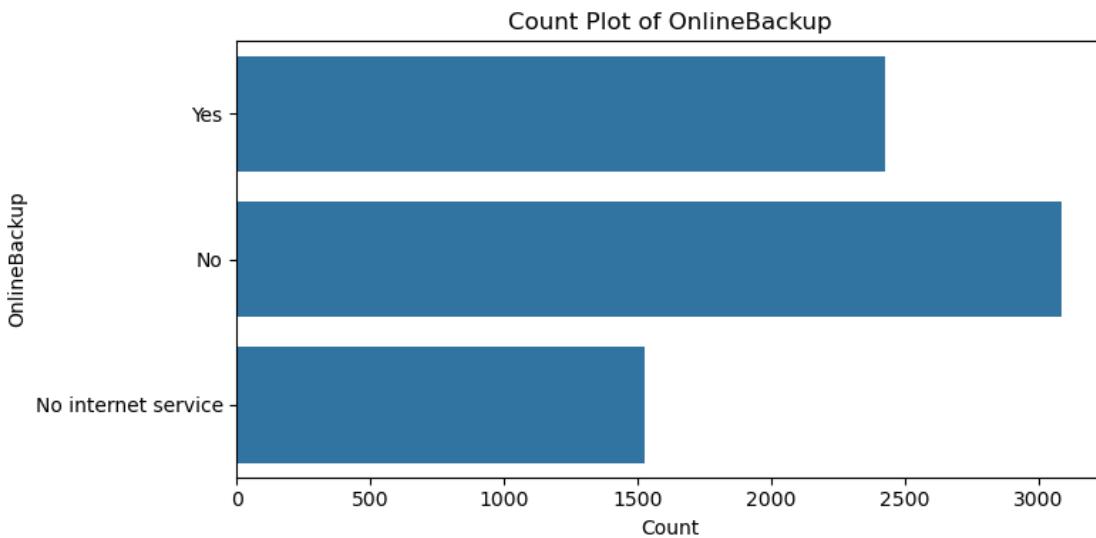
Count Plot of PhoneService

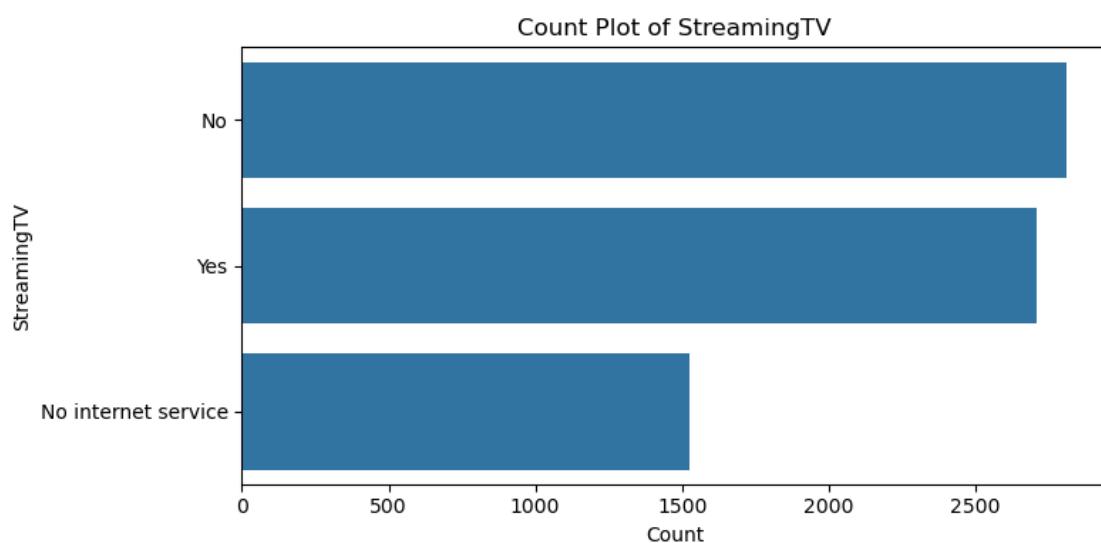
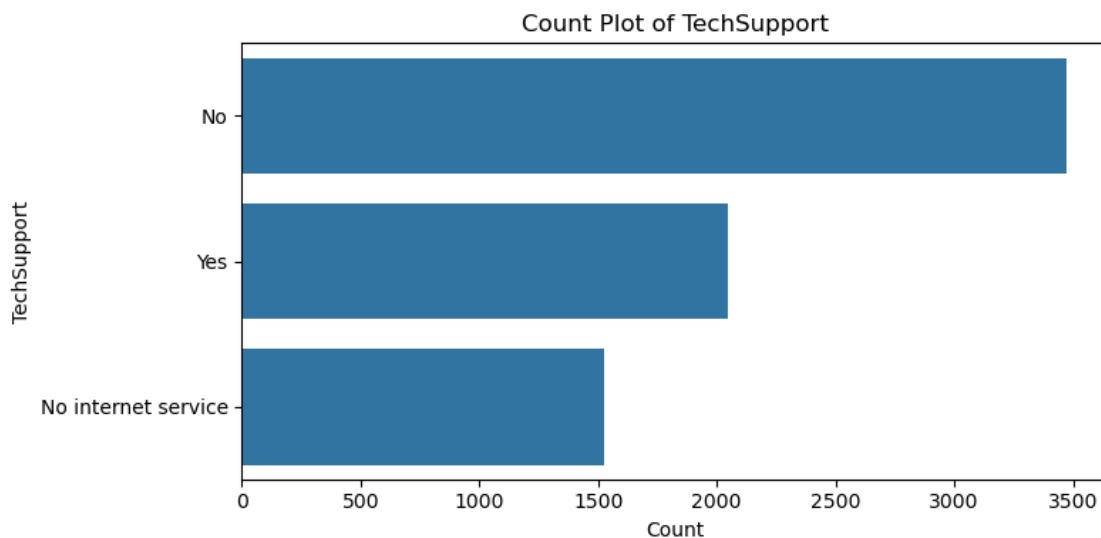


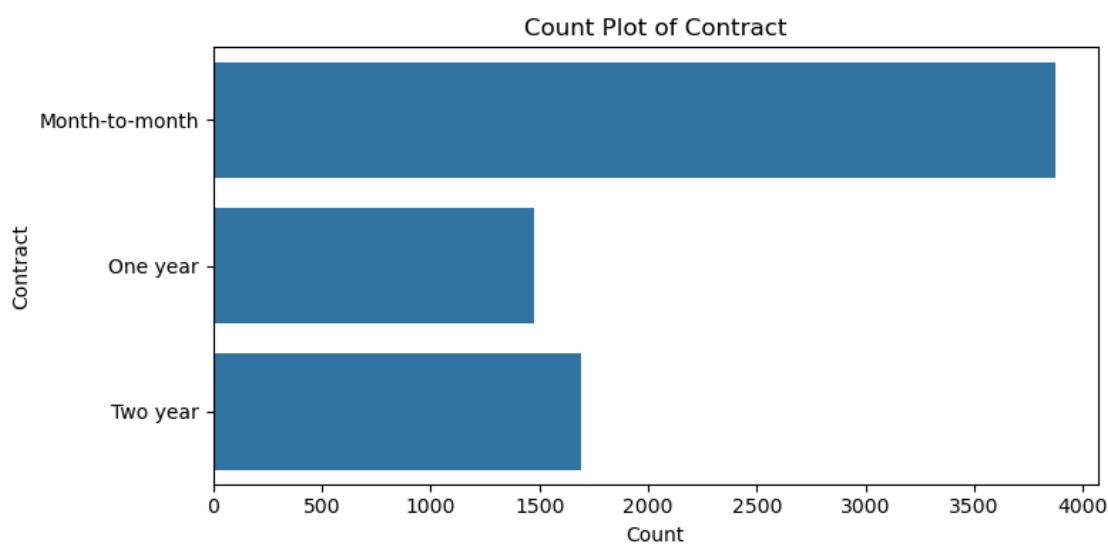
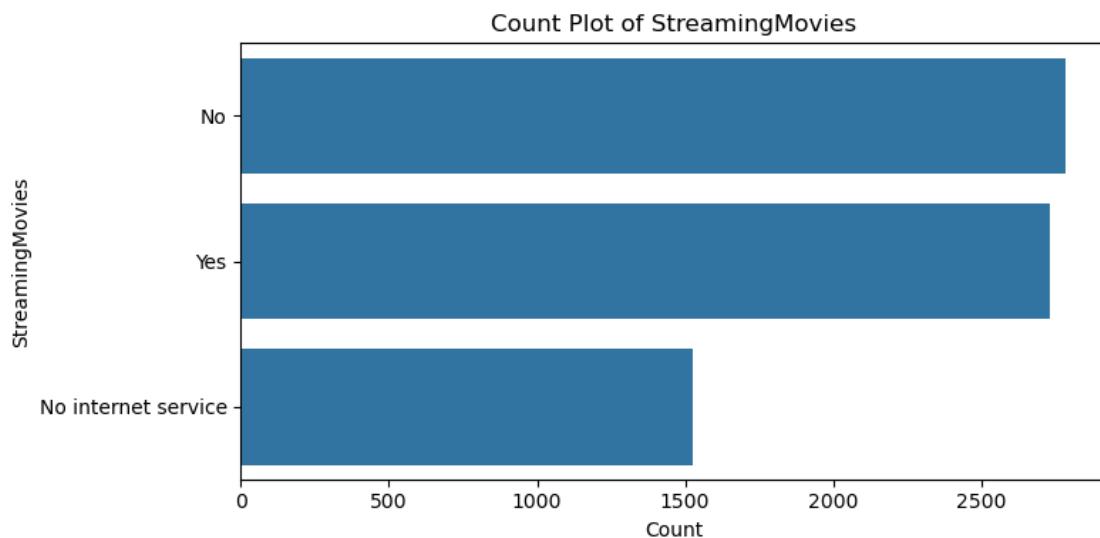
Count Plot of MultipleLines



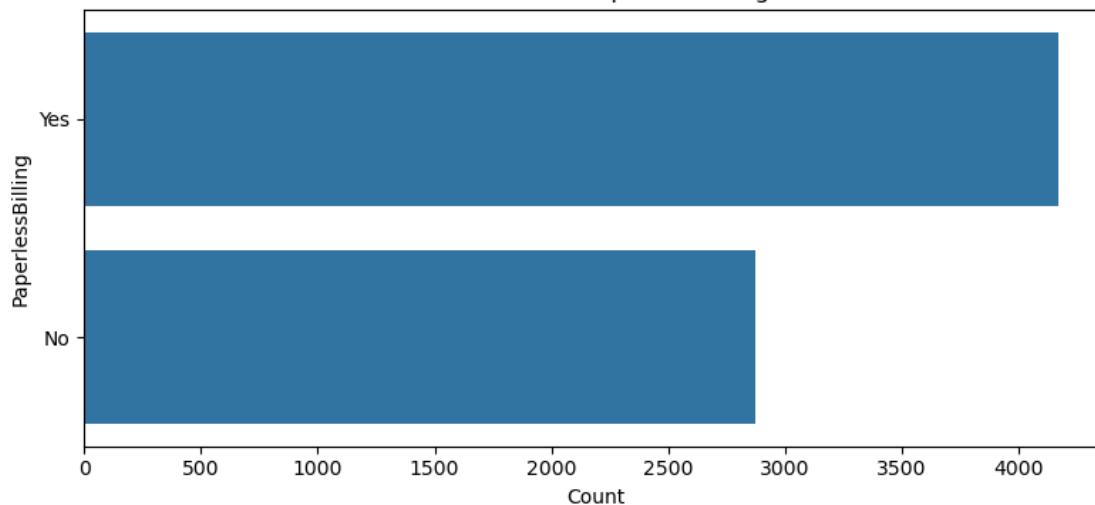




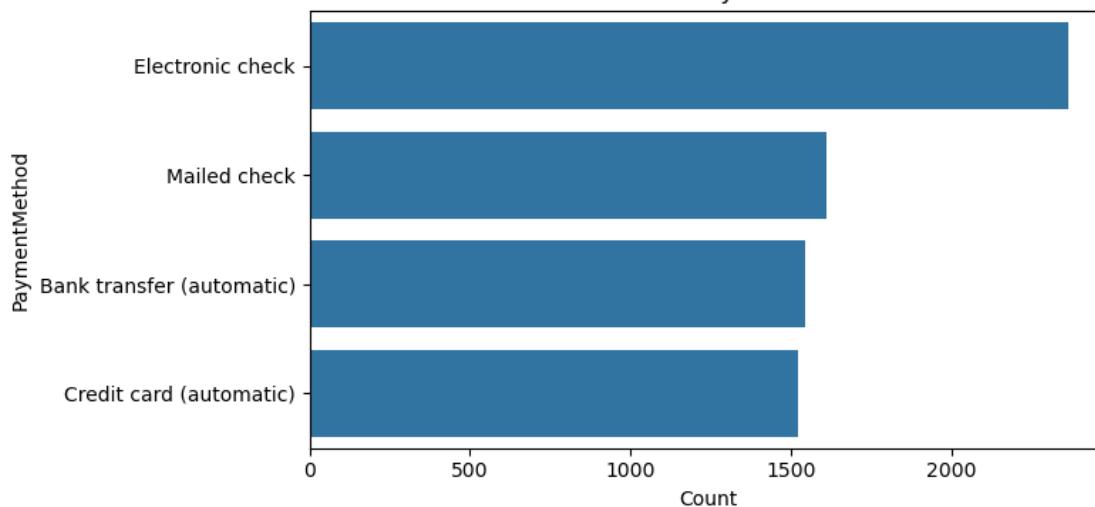


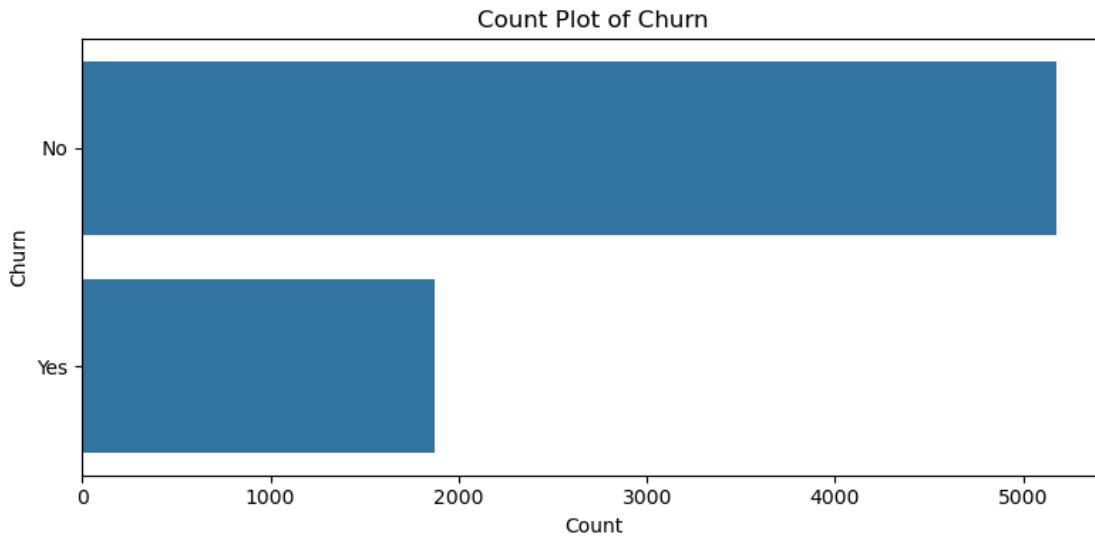


Count Plot of PaperlessBilling



Count Plot of PaymentMethod





#### 0.4 4. Data Preprocessing

```
[152]: df.head(3)
```

```
[152]: gender SeniorCitizen Partner Dependents tenure PhoneService \
0 Female 0 Yes No 1 No
1 Male 0 No No 34 Yes
2 Male 0 No No 2 Yes

MultipleLines InternetService OnlineSecurity OnlineBackup \
0 No phone service DSL No Yes
1 No DSL Yes No
2 No DSL Yes Yes

DeviceProtection TechSupport StreamingTV StreamingMovies Contract \
0 No No No No Month-to-month
1 Yes No No No One year
2 No No No No Month-to-month

PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn
0 Yes Electronic check 29.85 29.85 No
1 No Mailed check 56.95 1889.50 No
2 Yes Mailed check 53.85 108.15 Yes
```

#### 0.4.1 Label encoding of target column

```
[153]: df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
C:\Users\asus1\AppData\Local\Temp\ipykernel_27480\2364848822.py:1:
FutureWarning: Downcasting behavior in `replace` is deprecated and will be
removed in a future version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
[154]: df.head()
```

```
[154]:    gender SeniorCitizen Partner Dependents tenure PhoneService \
0   Female          0      Yes        No       1        No
1     Male          0       No        No      34       Yes
2     Male          0       No        No       2       Yes
3     Male          0       No        No      45        No
4   Female          0       No        No       2       Yes

           MultipleLines InternetService OnlineSecurity OnlineBackup \
0  No phone service             DSL         No       Yes
1            No                  DSL        Yes       No
2            No                  DSL        Yes       Yes
3  No phone service             DSL        Yes       No
4            No  Fiber optic         No       No

      DeviceProtection TechSupport StreamingTV StreamingMovies      Contract \
0            No          No        No        No Month-to-month
1           Yes          No        No        No    One year
2            No          No        No        No Month-to-month
3           Yes          Yes       No        No    One year
4            No          No        No        No Month-to-month

      PaperlessBilling      PaymentMethod MonthlyCharges  TotalCharges \
0            Yes  Electronic check      29.85        29.85
1            No    Mailed check      56.95     1889.50
2            Yes    Mailed check      53.85       108.15
3            No  Bank transfer (automatic)  42.30     1840.75
4            Yes  Electronic check      70.70       151.65

      Churn
0      0
1      0
2      1
3      0
4      1
```

```
[155]: print(df["Churn"].value_counts())
```

```
Churn
0    5174
1    1869
Name: count, dtype: int64
```

#### 0.4.2 Label encoding of categorical features

```
[156]: # identifying columns with object data type
object_columns = df.select_dtypes(include="object").columns
```

```
[157]: print(object_columns)
```

```
Index(['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
       'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
       'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
       'PaperlessBilling', 'PaymentMethod'],
      dtype='object')
```

```
[158]: # initialize a dictionary to save the encoders
encoders = {}
```

```
# apply label encoding and store the encoders
for column in object_columns:
    label_encoder = LabelEncoder()
    df[column] = label_encoder.fit_transform(df[column])
    encoders[column] = label_encoder

# save the encoders to a pickle file
with open("encoders.pkl", "wb") as f:
    pickle.dump(encoders, f)
```

```
[159]: encoders
```

```
[159]: {'gender': LabelEncoder(),
         'Partner': LabelEncoder(),
         'Dependents': LabelEncoder(),
         'PhoneService': LabelEncoder(),
         'MultipleLines': LabelEncoder(),
         'InternetService': LabelEncoder(),
         'OnlineSecurity': LabelEncoder(),
         'OnlineBackup': LabelEncoder(),
         'DeviceProtection': LabelEncoder(),
         'TechSupport': LabelEncoder(),
         'StreamingTV': LabelEncoder(),
```

```
'StreamingMovies': LabelEncoder(),
'Contract': LabelEncoder(),
'PaperlessBilling': LabelEncoder(),
'PaymentMethod': LabelEncoder()}
```

[160]: df.head()

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	
	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\		
0	1	0	0	2			
1	0	0	2	0			
2	0	0	2	2			
3	1	0	2	0			
4	0	1	0	0			
	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\	
0	0	0	0	0	0	0	
1	2	0	0	0	0	1	
2	0	0	0	0	0	0	
3	2	2	0	0	0	1	
4	0	0	0	0	0	0	
	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn		
0	1	2	29.85	29.85	0		
1	0	3	56.95	1889.50	0		
2	1	3	53.85	108.15	1		
3	0	0	42.30	1840.75	0		
4	1	2	70.70	151.65	1		

[161]: df.tail()

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
7038	1	0	1	1	24	1	
7039	0	0	1	1	72	1	
7040	0	0	1	1	11	0	
7041	1	1	1	0	4	1	
7042	1	0	0	0	66	1	
	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\		
7038	2	0	2	0			
7039	2	1	0	2			

7040	1	0	2	0
7041	2	1	0	0
7042	0	1	2	0
7038	2	2	2	2
7039	2	0	2	2
7040	0	0	0	0
7041	0	0	0	0
7042	2	2	2	2
7038	1	3	84.80	1990.50
7039	1	1	103.20	7362.90
7040	1	2	29.60	346.45
7041	1	3	74.40	306.60
7042	1	0	105.65	6844.50

#### 0.4.3 Traianing and test data split

```
[162]: # splitting the features and target
X = df.drop(columns=["Churn"])
y = df["Churn"]
```

```
[163]: # split training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
[164]: print(y_train.shape)
```

(5634,)

```
[165]: print(y_train.value_counts())
```

```
Churn
0    4138
1    1496
Name: count, dtype: int64
```

#### 0.4.4 Synthetic Minority Oversampling TECnique (SMOTE)

```
[166]: smote = SMOTE(random_state=42)
```

```
[167]: X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
[168]: print(y_train_smote.shape)
```

(8276,)

```
[169]: print(y_train_smote.value_counts())
```

```
Churn
0    4138
1    4138
Name: count, dtype: int64
```

## 0.5 5. Model Training

### 0.5.1 Training with default hyperparameters

```
[170]: # dictionary of models
models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "XGBoost": XGBClassifier(random_state=42)
}
```

```
[171]: # dictionary to store the cross validation results
cv_scores = {}

# perform 5-fold cross validation for each model
for model_name, model in models.items():
    print(f"Training {model_name} with default parameters")
    scores = cross_val_score(model, X_train_smote, y_train_smote, cv=5,
                             scoring="accuracy")
    cv_scores[model_name] = scores
    print(f"{model_name} cross-validation accuracy: {np.mean(scores):.2f}")
    print("-" * 70)
```

Training Decision Tree with default parameters

Decision Tree cross-validation accuracy: 0.78

---

Training Random Forest with default parameters

Random Forest cross-validation accuracy: 0.84

---

Training XGBoost with default parameters

XGBoost cross-validation accuracy: 0.83

---

```
[172]: cv_scores
```

```
[172]: {'Decision Tree': array([0.68115942, 0.71903323, 0.81752266, 0.84350453,
 0.84350453]),
 'Random Forest': array([0.72705314, 0.76676737, 0.90453172, 0.89244713,
 0.89848943]),
```

```
'XGBoost': array([0.71074879, 0.75226586, 0.90271903, 0.89123867, 0.89909366])}
```

### 0.5.2 Random Forest gives the highest accuracy compared to other models with default parameters

```
[173]: rfc = RandomForestClassifier(random_state=42)
```

```
[174]: rfc.fit(X_train_smote, y_train_smote)
```

```
[174]: RandomForestClassifier(random_state=42)
```

```
[175]: print(y_test.value_counts())
```

```
Churn
0    1036
1     373
Name: count, dtype: int64
```

## 0.6 6. Model Evaluation

```
[176]: # evaluate on test data
y_test_pred = rfc.predict(X_test)

print("Accuracy Score:\n", accuracy_score(y_test, y_test_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred))
print("Classification Report:\n", classification_report(y_test, y_test_pred))
```

```
Accuracy Score:
0.7771469127040455
Confusion Matrix:
[[880 156]
 [158 215]]
Classification Report:
      precision    recall  f1-score   support
          0       0.85     0.85     0.85      1036
          1       0.58     0.58     0.58      373

   accuracy                           0.78      1409
  macro avg       0.71     0.71     0.71      1409
weighted avg       0.78     0.78     0.78      1409
```

```
[177]: # save the trained model as a pickle file
model_data = {"model": rfc, "features_names": X.columns.tolist()}
```

```
with open("customer_churn_model.pkl", "wb") as f:  
    pickle.dump(model_data, f)
```

## 0.7 7. Load the saved model and build a Predictive System

```
[178]: # load teh saved model and the feature names
```

```
with open("customer_churn_model.pkl", "rb") as f:  
    model_data = pickle.load(f)  
  
loaded_model = model_data["model"]  
feature_names = model_data["features_names"]
```

```
[179]: print(loaded_model)
```

```
RandomForestClassifier(random_state=42)
```

```
[180]: print(feature_names)
```

```
['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService',  
'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',  
'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges']
```

```
[181]: input_data = {  
    'gender': 'Female',  
    'SeniorCitizen': 0,  
    'Partner': 'Yes',  
    'Dependents': 'No',  
    'tenure': 1,  
    'PhoneService': 'No',  
    'MultipleLines': 'No phone service',  
    'InternetService': 'DSL',  
    'OnlineSecurity': 'No',  
    'OnlineBackup': 'Yes',  
    'DeviceProtection': 'No',  
    'TechSupport': 'No',  
    'StreamingTV': 'No',  
    'StreamingMovies': 'No',  
    'Contract': 'Month-to-month',  
    'PaperlessBilling': 'Yes',  
    'PaymentMethod': 'Electronic check',  
    'MonthlyCharges': 29.85,  
    'TotalCharges': 29.85  
}
```

```

input_data_df = pd.DataFrame([input_data])

with open("encoders.pkl", "rb") as f:
    encoders = pickle.load(f)

# encode categorical features using teh saved encoders
for column, encoder in encoders.items():
    input_data_df[column] = encoder.transform(input_data_df[column])

# make a prediction
prediction = loaded_model.predict(input_data_df)
pred_prob = loaded_model.predict_proba(input_data_df)

print(prediction)

# results
print(f"Prediction: {'Churn' if prediction[0] == 1 else 'No Churn'}")
print(f"Prediciton Probability: {pred_prob}")

```

[0]  
Prediction: No Churn  
Prediciton Probability: [[0.83 0.17]]

## 0.8 Conclusion

This project demonstrates an end-to-end machine learning workflow for predicting customer churn.

Through effective data analysis, preprocessing, and model building, meaningful churn patterns were identified.

The final model provides actionable insights that can help businesses improve customer retention.

Overall, this project reflects practical application of machine learning in a real-world business scenario.

[ ]: